

nce-capstone-project-health-care

October 7, 2023

0.1 Data Science Capstone Project : Health Care

```
[1]: # Importing the libraries

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
%matplotlib inline
from matplotlib import style
import seaborn as sns
import imblearn
```

```
[2]: # Loading the dataset

df = pd.read_csv('health care diabetes.csv')
```

0.1.1 Data Exploration:

```
[3]: df.head() #top 5 dataset
```

```
[3]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	\
0	6	148	72	35	0	33.6	
1	1	85	66	29	0	26.6	
2	8	183	64	0	0	23.3	
3	1	89	66	23	94	28.1	
4	0	137	40	35	168	43.1	

	DiabetesPedigreeFunction	Age	Outcome
0	0.627	50	1
1	0.351	31	0
2	0.672	32	1
3	0.167	21	0
4	2.288	33	1

```
[4]: df.tail() #bottom 5 dataset
```

```
[4]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI \
763	10	101	76	48	180	32.9
764	2	122	70	27	0	36.8
765	5	121	72	23	112	26.2
766	1	126	60	0	0	30.1
767	1	93	70	31	0	30.4

	DiabetesPedigreeFunction	Age	Outcome
763	0.171	63	0
764	0.340	27	0
765	0.245	30	0
766	0.349	47	1
767	0.315	23	0

```
[5]: df.info() #concise information about a DataFrame.
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 768 entries, 0 to 767
Data columns (total 9 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Pregnancies           768 non-null   int64
1   Glucose               768 non-null   int64
2   BloodPressure         768 non-null   int64
3   SkinThickness         768 non-null   int64
4   Insulin               768 non-null   int64
5   BMI                   768 non-null   float64
6   DiabetesPedigreeFunction 768 non-null   float64
7   Age                   768 non-null   int64
8   Outcome               768 non-null   int64
dtypes: float64(2), int64(7)
memory usage: 54.1 KB
```

```
[6]: df.shape #Rows and columns
```

```
[6]: (768, 9)
```

```
[7]: df.describe() # statistical summary information for each numeric column
```

```
[7]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin \
count	768.000000	768.000000	768.000000	768.000000	768.000000
mean	3.845052	120.894531	69.105469	20.536458	79.799479
std	3.369578	31.972618	19.355807	15.952218	115.244002
min	0.000000	0.000000	0.000000	0.000000	0.000000
25%	1.000000	99.000000	62.000000	0.000000	0.000000
50%	3.000000	117.000000	72.000000	23.000000	30.500000
75%	6.000000	140.250000	80.000000	32.000000	127.250000

max	17.000000	199.000000	122.000000	99.000000	846.000000
-----	-----------	------------	------------	-----------	------------

	BMI	DiabetesPedigreeFunction	Age	Outcome
count	768.000000	768.000000	768.000000	768.000000
mean	31.992578	0.471876	33.240885	0.348958
std	7.884160	0.331329	11.760232	0.476951
min	0.000000	0.078000	21.000000	0.000000
25%	27.300000	0.243750	24.000000	0.000000
50%	32.000000	0.372500	29.000000	0.000000
75%	36.600000	0.626250	41.000000	1.000000
max	67.100000	2.420000	81.000000	1.000000

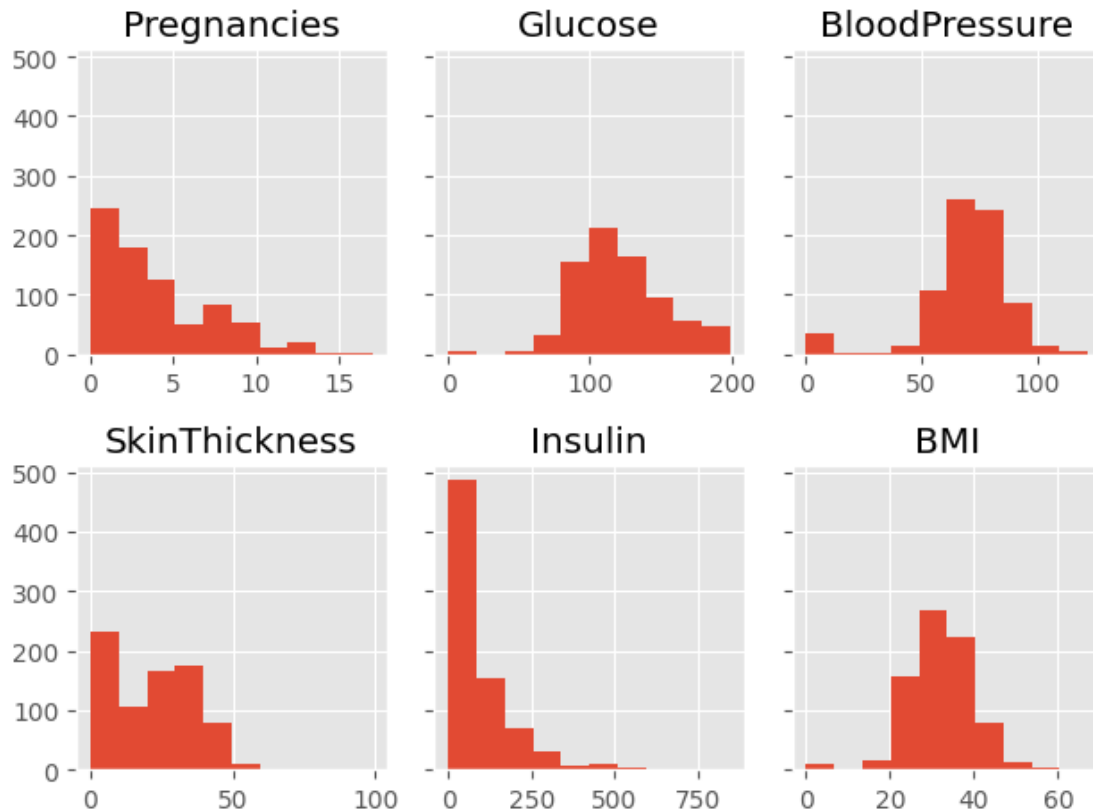
- As per descriptive analysis we see that minimum value is 0 for columns such as **Pregnancies**, **Glucose**, **SkinThickness**, **Insulin**, **BMI** etc.
- which is practically impossible so we have to treat this as null value (**NAN**) hence replace with its **median** values.

```
[8]: ## Visualizations
style.use('ggplot')
fig,axs=plt.subplots(2,3,sharey=True,tight_layout=True)

axs[0,0].hist(df['Pregnancies'],bins=10)
axs[0,1].hist(df['Glucose'],bins=10)
axs[0,2].hist(df['BloodPressure'],bins=10)
axs[1,0].hist(df['SkinThickness'],bins=10)
axs[1,1].hist(df['Insulin'],bins=10)
axs[1,2].hist(df['BMI'],bins=10)

#set title
axs[0,0].set_title('Pregnancies')
axs[0,1].set_title('Glucose')
axs[0,2].set_title('BloodPressure')
axs[1,0].set_title('SkinThickness')
axs[1,1].set_title('Insulin')
axs[1,2].set_title('BMI')
```

```
[8]: Text(0.5, 1.0, 'BMI')
```



```
[9]: # Treating the missing values (NaN) with the meidan(50%)

df['Pregnancies'] = df['Pregnancies'].replace(0, 3.000000)
df['Glucose'] = df['Glucose'].replace(0, 117.000000)
df['BloodPressure'] = df['BloodPressure'].replace(0, 72.000000)
df['SkinThickness'] = df['SkinThickness'].replace(0, 23.000000)
df['Insulin'] = df['Insulin'].replace(0, 30.500000)
df['BMI'] = df['BMI'].replace(0, 32.000000)
```

```
[10]: df.describe()
```

```
[10]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin \
count	768.000000	768.000000	768.000000	768.000000	768.000000
mean	4.278646	121.656250	72.386719	27.334635	94.652344
std	3.021516	30.438286	12.096642	9.229014	105.547598
min	1.000000	44.000000	24.000000	7.000000	14.000000
25%	2.000000	99.750000	64.000000	23.000000	30.500000
50%	3.000000	117.000000	72.000000	23.000000	31.250000
75%	6.000000	140.250000	80.000000	32.000000	127.250000
max	17.000000	199.000000	122.000000	99.000000	846.000000


	BMI	DiabetesPedigreeFunction	Age	Outcome
count	768.000000	768.000000	768.000000	768.000000
mean	32.450911	0.471876	33.240885	0.348958
std	6.875366	0.331329	11.760232	0.476951
min	18.200000	0.078000	21.000000	0.000000
25%	27.500000	0.243750	24.000000	0.000000
50%	32.000000	0.372500	29.000000	0.000000
75%	36.600000	0.626250	41.000000	1.000000
max	67.100000	2.420000	81.000000	1.000000

```
[11]: ## Visualizations
style.use('ggplot')
fig,axs=plt.subplots(2,3,sharey=True,tight_layout=True)
for ax in axs.flat:
    ax.set_facecolor('yellow')

axs[0,0].hist(df['Pregnancies'],bins=10)
axs[0,1].hist(df['Glucose'],bins=10)
axs[0,2].hist(df['BloodPressure'],bins=10)
axs[1,0].hist(df['SkinThickness'],bins=10)
axs[1,1].hist(df['Insulin'],bins=10)
axs[1,2].hist(df['BMI'],bins=10)

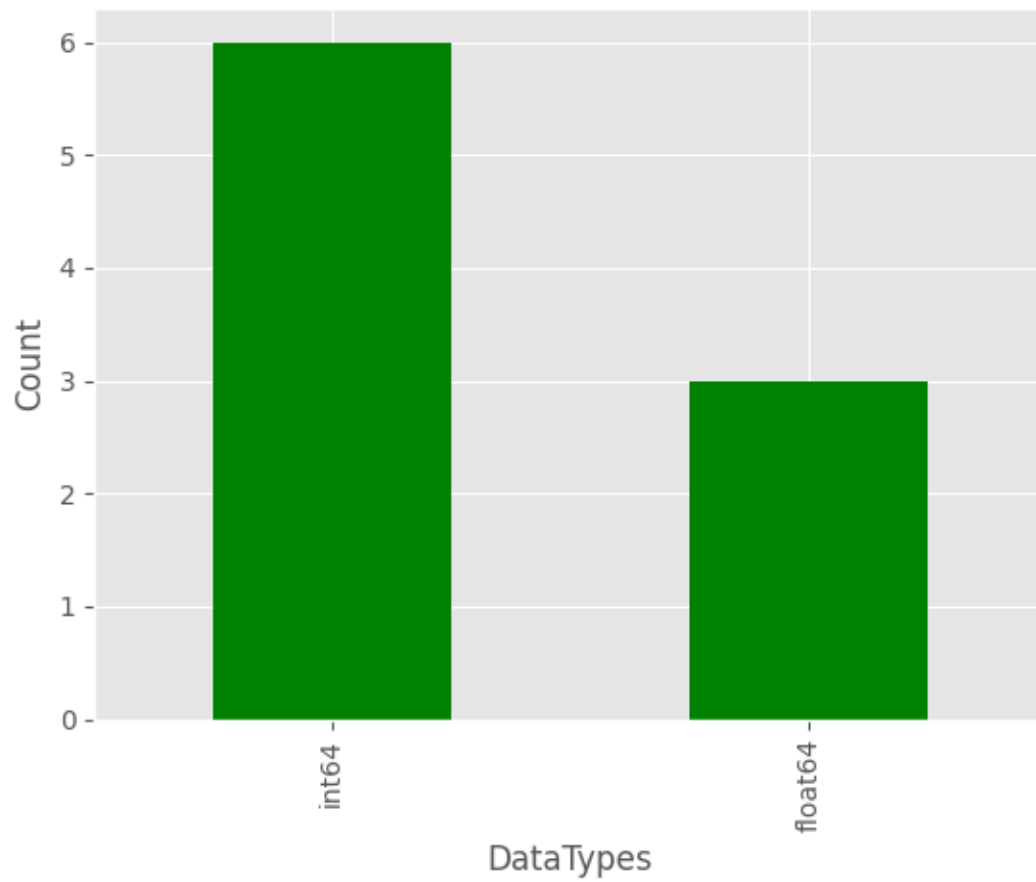
#set title
axs[0,0].set_title('Pregnancies')
axs[0,1].set_title('Glucose')
axs[0,2].set_title('BloodPressure')
axs[1,0].set_title('SkinThickness')
axs[1,1].set_title('Insulin')
axs[1,2].set_title('BMI')
```

```
[11]: Text(0.5, 1.0, 'BMI')
```



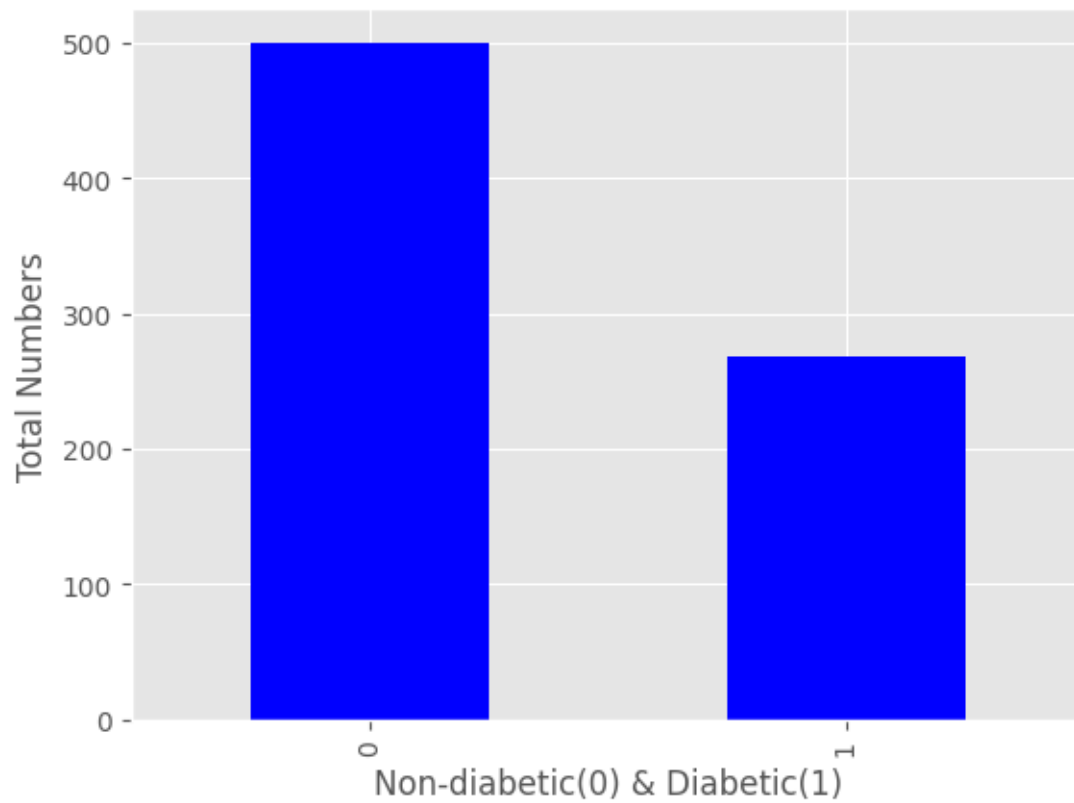
```
[12]: # Creating count(frequency) plot describing the data types and the count of
      ↪ variables.
      style.use('ggplot')
      df.dtypes.value_counts().plot(kind='bar', color='g')
      plt.xlabel('DataTypes')
      plt.ylabel('Count')
```

```
[12]: Text(0, 0.5, 'Count')
```



```
[13]: df['Outcome'].value_counts().plot(kind='bar', color='b')  
plt.xlabel('Non-diabetic(0) & Diabetic(1)')  
plt.ylabel('Total Numbers')
```

```
[13]: Text(0, 0.5, 'Total Numbers')
```

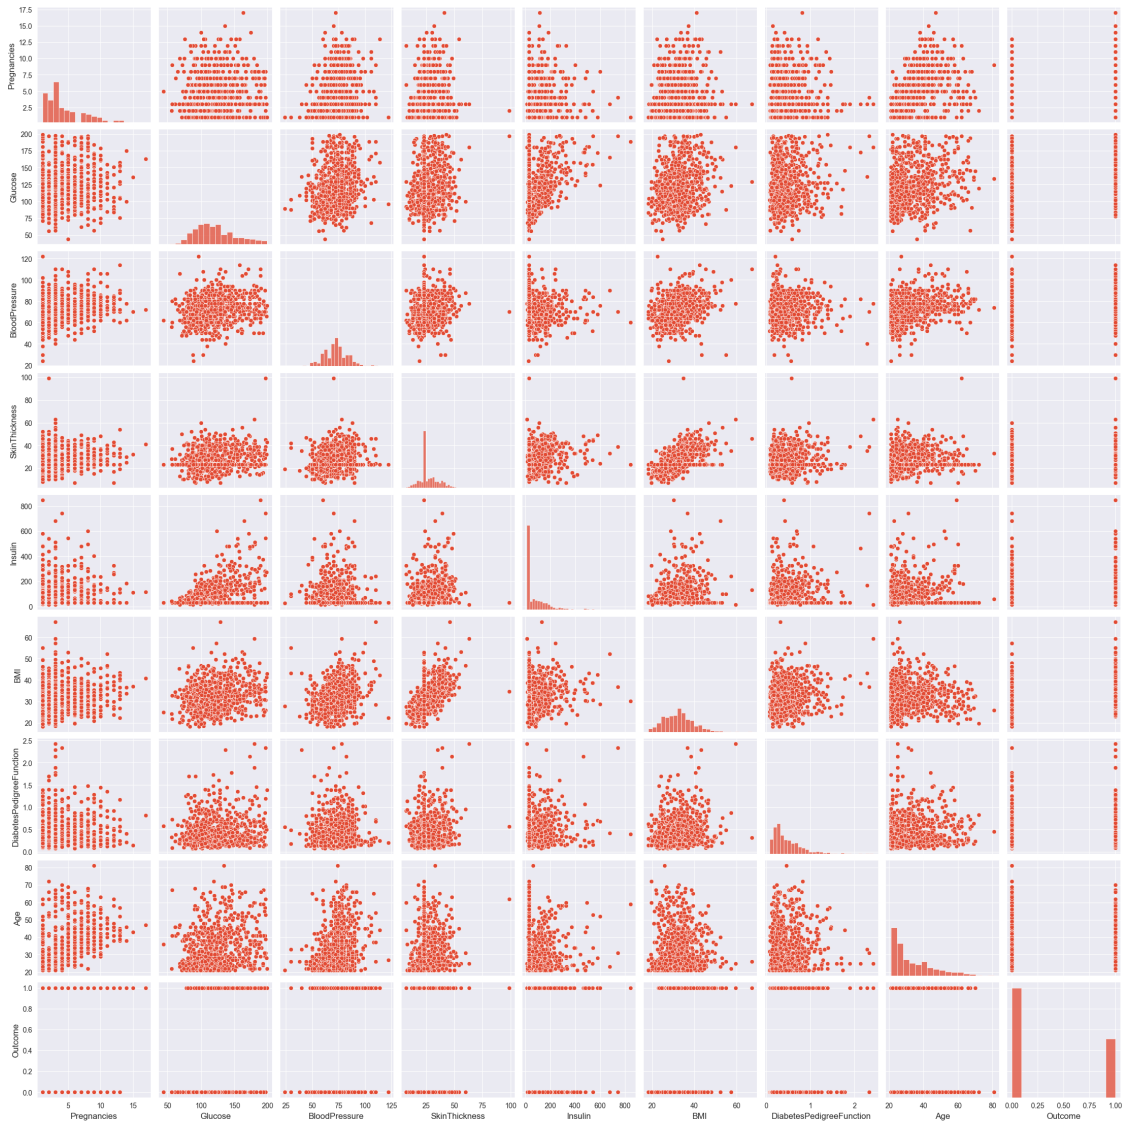


As we can see that there is an **imbalance** in data that means we might need to apply **SMOTE** moving forward.

```
[14]: sns.set_style('darkgrid')
      sns.pairplot(df)
```

```
C:\Users\Deepak Yadav\anaconda3\lib\site-packages\seaborn\axisgrid.py:118:
UserWarning: The figure layout has changed to tight
  self._figure.tight_layout(*args, **kwargs)
```

```
[14]: <seaborn.axisgrid.PairGrid at 0x1df618ca050>
```

```
[15]: df1=df.corr()
df1
```

```
[15]:
```


	Pregnancies	Glucose	BloodPressure	SkinThickness	\
Pregnancies	1.000000	0.149326	0.247530	0.060706	
Glucose	0.149326	1.000000	0.218937	0.172143	
BloodPressure	0.247530	0.218937	1.000000	0.147809	
SkinThickness	0.060706	0.172143	0.147809	1.000000	
Insulin	-0.059580	0.357573	-0.028721	0.238188	
BMI	0.080540	0.231400	0.281132	0.546951	
DiabetesPedigreeFunction	-0.016151	0.137327	-0.002378	0.142977	
Age	0.538169	0.266909	0.324915	0.054514	
Outcome	0.245466	0.492782	0.165723	0.189065	

	Insulin	BMI	DiabetesPedigreeFunction \
Pregnancies	-0.059580	0.080540	-0.016151
Glucose	0.357573	0.231400	0.137327
BloodPressure	-0.028721	0.281132	-0.002378
SkinThickness	0.238188	0.546951	0.142977
Insulin	1.000000	0.189022	0.178029
BMI	0.189022	1.000000	0.153506
DiabetesPedigreeFunction	0.178029	0.153506	1.000000
Age	-0.015413	0.025744	0.033561
Outcome	0.148457	0.312249	0.173844

	Age	Outcome
Pregnancies	0.538169	0.245466
Glucose	0.266909	0.492782
BloodPressure	0.324915	0.165723
SkinThickness	0.054514	0.189065
Insulin	-0.015413	0.148457
BMI	0.025744	0.312249
DiabetesPedigreeFunction	0.033561	0.173844
Age	1.000000	0.238356
Outcome	0.238356	1.000000

```
[16]: sns.heatmap(df.corr(),annot=True) # Plotting the heatmap
```

```
[16]: <Axes: >
```



```
[17]: df1.to_csv('df_corr.csv')
```

```
[18]: # Seprating 'Independent' and 'Dependent' Variables.
```

```
x=df.drop(['Outcome'],axis=1)
y=df['Outcome']
```

```
[19]: # feature scaling
```

```
from sklearn.preprocessing import StandardScaler
```

```
[20]: scale = StandardScaler() #Initialize the StandardScaler
x=scale.fit_transform(x)
```

Balancing dataset using **SMOTE**

```
[21]: from imblearn.over_sampling import SMOTE
sm = SMOTE(random_state=42)
```

```
[22]: X,Y=sm.fit_resample(x,y)
```

```
[23]: X.shape, Y.shape
```

```
[23]: ((1000, 8), (1000,))
```

0.1.2 Data Modeling:

```
[24]: #train test split

from sklearn.model_selection import train_test_split
X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size=0.25,
↳random_state=42)
```

Building the **Machine Learning pipeline** for different ML Classification Models like **KNN**, **logistic Regression**, **DecisionTree**, **RandomForest**, **SVM** & **NB** inorder to test multiple algorithms to see which one works best for this specific problem.

```
[25]: from sklearn.neighbors import KNeighborsClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.svm import SVC
from sklearn.naive_bayes import GaussianNB
```

Creating a list called **model_pipeline** and appending various machine learning classifiers to it.

```
[26]: model_pipeline=[]
model_pipeline.append(KNeighborsClassifier())
model_pipeline.append(LogisticRegression())
model_pipeline.append(DecisionTreeClassifier())
model_pipeline.append(RandomForestClassifier())
model_pipeline.append(SVC())
model_pipeline.append(GaussianNB())
```

Importing several important **metrics and functions** from scikit-learn for **evaluating machine learning models**.

```
[27]: from sklearn.metrics import
↳accuracy_score,classification_report,confusion_matrix,roc_curve,auc
```

Creating a list of machine learning models (**model_list**) and using a list of classifiers (**model_pipeline**) to train and assess them. Evaluating metrics like **test accuracy**, **training accuracy**, **AUC**, and **confusion matrices** for each model.

```
[28]: model_list = ['Logistic Regression', 'Decision Tree', 'Random Forest', 'SVM',
↳'KNN', 'GaussianNB']
test_acc_list = []
```

```

train_acc_list = []
auc_list = []
cm_list = []

for model in model_pipeline:
    model.fit(X_train, Y_train)
    Y_pred = model.predict(X_test)
    Y_train_pred = model.predict(X_train)
    test_acc_list.append(accuracy_score(Y_test, Y_pred))
    train_acc_list.append(accuracy_score(Y_train, Y_train_pred))
    fpr, tpr, _thresholds = roc_curve(Y_test, Y_pred)
    auc_list.append(round(auc(fpr, tpr), 2))
    cm_list.append(confusion_matrix(Y_test, Y_pred))

```

Creating a DataFrame named **result_df** containing various metrics for different machine learning models. The DataFrame includes columns for the model names, test accuracy, training accuracy, and AUC (Area Under the ROC Curve).

```
[29]: result_df=pd.DataFrame({'model':model_list,'test_accuracy':
    ↪test_acc_list,'train_accuracy':train_acc_list,'AUC':auc_list})
```

```
[30]: result_df
```

```
[30]:
```

	model	test_accuracy	train_accuracy	AUC
0	Logistic Regression	0.756	0.836000	0.75
1	Decision Tree	0.752	0.761333	0.75
2	Random Forest	0.736	1.000000	0.74
3	SVM	0.784	1.000000	0.78
4	KNN	0.784	0.860000	0.78
5	GaussianNB	0.748	0.722667	0.75

We observe an instance of **overfitting** in the data, which prompts us to consider **cross-validation**. In this case, we opt for **K-Fold cross-validation**.

In this scenario, it's worth noting that **K-Nearest Neighbors (KNN)** exhibits lower accuracy compared to **Support Vector Machines (SVM)** and **Random Forest**. As a result, we decide to implement **K-Fold cross-validation** for the **Random Forest model**.

```
[31]: from sklearn.ensemble import AdaBoostClassifier
base_model = RandomForestClassifier(max_depth=2)
model = AdaBoostClassifier(base_model, n_estimators=50)
model.fit(X_train, Y_train)
Y_pred = model.predict(X_test)
accuracy = accuracy_score(Y_test, Y_pred)
conf_matrix = confusion_matrix(Y_test, Y_pred)
class_report = classification_report(Y_test, Y_pred)
data = {
    'Metric': ['Accuracy', 'Confusion Matrix', 'Classification Report'],
```

```

    'Result': [accuracy, conf_matrix, class_report]
}
results_df = pd.DataFrame(data)
print(results_df)

```

	Metric	Result
0	Accuracy	0.744
1	Confusion Matrix	[[89, 35], [29, 97]]
2	Classification Report	precision recall f1-score ...

```
[32]: print(classification_report(Y_pred,Y_test))
```

	precision	recall	f1-score	support
0	0.72	0.75	0.74	118
1	0.77	0.73	0.75	132
accuracy			0.74	250
macro avg	0.74	0.74	0.74	250
weighted avg	0.75	0.74	0.74	250

Using the **XGBoost** library (xgboost) to build a binary classification model.

```
[33]: import xgboost as xgb
dtrain=xgb.DMatrix(X_train,label=Y_train)
params={'objective':'binary:logistic',
        'eval_metric':'logloss',
        'max_depth':3,
        'learning_rate':0.1,
        'gamma':0.2,
        'subsample':0.8,
        'colsample_bytree':0.8}
```

```
[34]: model=xgb.train(params,dtrain,num_boost_round=100)
```

```
[35]: dtest=xgb.DMatrix(X_test)
```

```
[36]: Y_pred_proba=model.predict(dtest)
```

```
[37]: Y_pred=[1 if p>=0.5 else 0 for p in Y_pred_proba]
```

```
[38]: acc=accuracy_score(Y_test,Y_pred)
```

```
[39]: acc
```

```
[39]: 0.76
```

In summary, we have determined that **SVM** exhibits **superior accuracy** and **AUC**. Additionally, SVM stands out as the appropriate model for this dataset due to its **absence of overfitting issues**.

[]: